

BigData: Apriori & SONS

A comparison among several implementations of Association-Rule based algorithms.

Details

Project

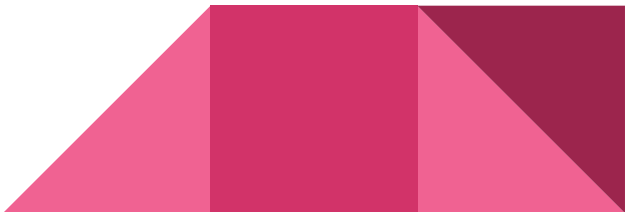
The project have been developed as framework for benchmarking a series of frequent itemset algorithms against various datasets and minimum support levels. All the experiments and results have been obtained by trying as much as possible to replicate equal testing conditions. Loading of data, running of the experiments, profiling and result storage are unified and are not included in the computation of the timings, just the pure execution of the algorithm.

Hardware

- Ubuntu 22.04
- 32GB RAM
- AMD Ryzen 9 5900X
- 12 Cores
- 24 Threads (*24 workers*)
- Spark 3.4.0 for Hadoop 3.3.4
- Python 3.10.9 (conda)

Algorithms

All implementations are based on PySpark and make use map, reduce, cartesian, filter and other typical paradigms from functional programming, replacing loops as much as possible. Some implementations, like `apriori_list` are made in pure python as they are intended to be used on partition data, which by nature in Spark, is not an RDD, but simple python object. Also in this case the same guideline have been followed.

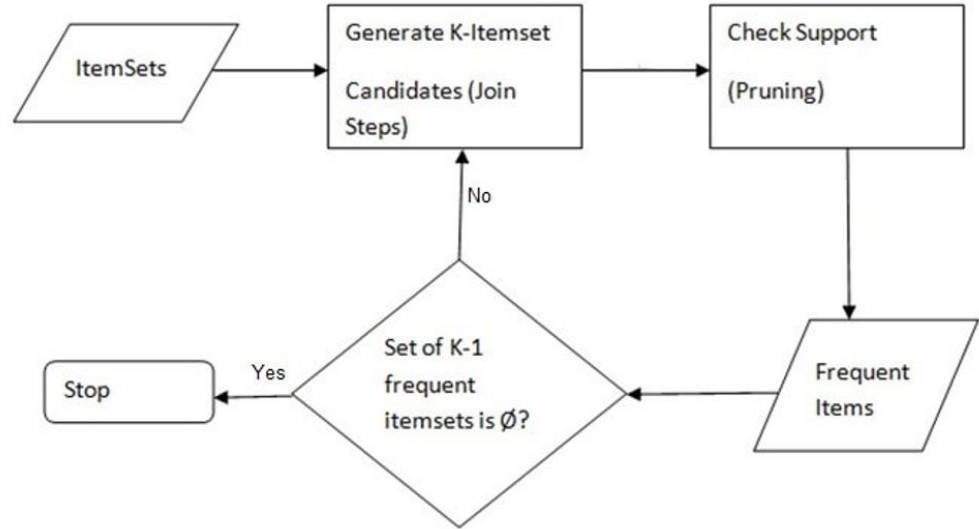


ALGORITHMS

Apriori

Steps

1. Count single transaction items
2. Obtain FrequentSet of size 1
3. Permutate to obtain candidates size 2
4. Perform cartesian between candidates and transactions.
5. Filter only existing sets for existing transactions and count them.
6. Obtain FrequentSet on size 2
7. Repeat from step 3



Apriori: Focus

Extending the Set

```
extendedSets = product(itemSupport, itemUnique) ★  
extendedSets = map(flatMergeCartesian, extendedSets) # ( ('3','1'),'4') -> ('1','3','4')  
extendedSets = tuple(filter(lambda item: len(item) == set_length, extendedSets))  
# from flatMergeCartesian: ( ('4','3'), '3') -> ('3','4')  
extendedSets = tuple(set(extendedSets))
```

Filtering the Set

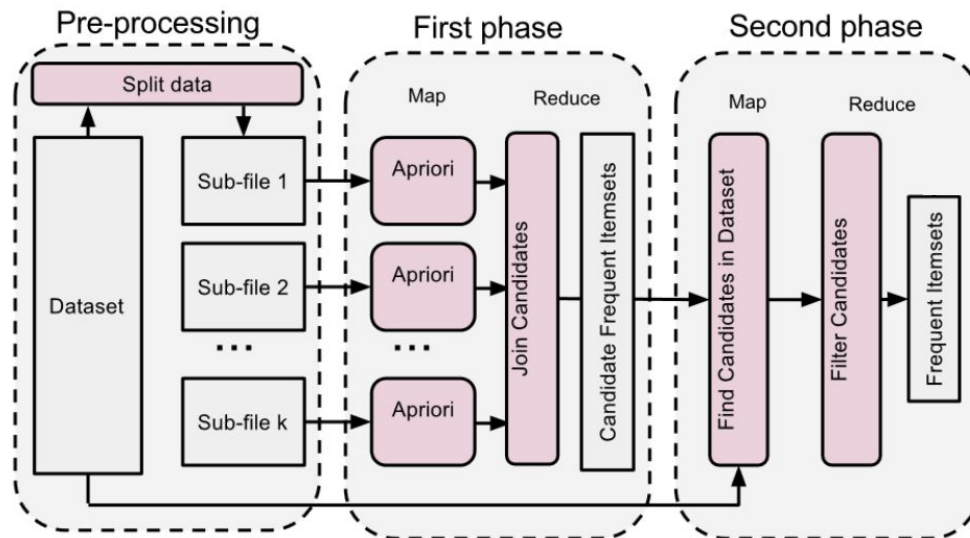
```
existingSets = product(extendedSets, transactions) ★  
existingSets = tuple(filter(lambda item: set(item[0]).issubset(item[1]), existingSets))  
existingSets = tuple(map(lambda item: item[0], existingSets))
```



SON

Steps

1. Partition the dataset
2. Perform “apriori” for each partition, discard the counting.
3. Merge the results from all partitions and use those as candidates.
4. Perform cartesian between candidates and transactions.
5. Find candidate sets.
6. Filter candidates and obtain Frequent Itemsets.



SON: Focus

Find candidates from partitions

```
candidateSets = partitionSets.mapPartitions(lambda part: ★priori_list(list(part), min_support))
candidateSets = candidateSets.map(lambda item: item[0])
candidateSets = candidateSets.coalesce(1)
candidateSets = candidateSets.distinct()
```

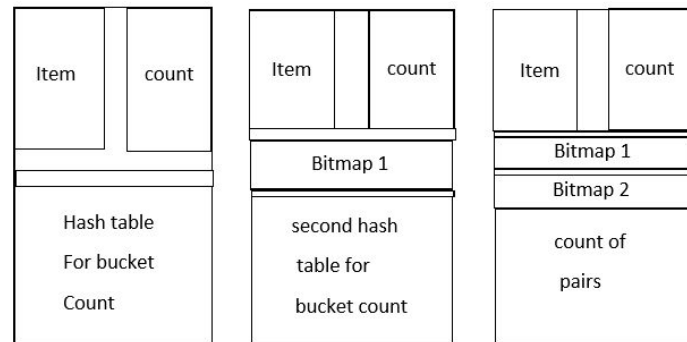
Test candidates

```
# test candidates
testingSets = transactions.cartesian(candidateSets) ★
testingSets = testingSets.filter(lambda item: set(item[0]).issubset(item[1]))
testingSets = testingSets.reduceByKey(lambda x, y: x + y)
testingSets = testingSets.filter(lambda item: item[1] > candidate_frequency)
```

PCY

Steps

1. Count single transaction items
2. Obtain FrequentSet of size 1
3. Partition the dataset into buckets
4. Generate Candidate set from row combination.
5. Using a hash function count items per bucket and extract candidates and use it for filtering.
6. Perform cartesian between candidates and transactions.
7. Find candidate sets.
8. Filter candidates and obtain Frequent Itemsets.



Multistage Memory Mapping.

PCY: Key ideas

Each row contains, in potential, all the frequent/candidate itemsets.

Make efficient use of memory it's important, use hash tables to select candidates.

Frequent Itemsets of size N , are a combination, without repetition of the individual items of $N-1$.

Why PCY?

?

Each row can be exploded in its subsets: processing can be parallelized at row level.

Why buckets ?

?

You can sum and obtain final frequency.

Why multiply for hashing, to generate just candidates?

?

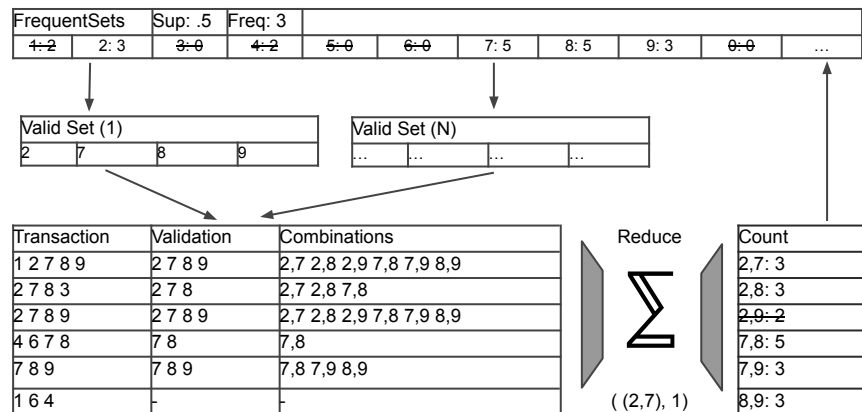
Remove every item you can before join and combinations.

Why joining huge tables?

PCY Bucketless

Steps

1. Count single transaction items
2. Obtain FrequentSets of size N=1.
3. Make the valid set, the FrequentSets on size N.
4. For each valid transaction line generate a list of actual sets of size N from the combination items of that row.
5. Each transaction line generates multiple sets, make them flat.
6. Reduce by key and filter.
7. Obtain Frequent Itemsets.



PCY Bucketless: Focus

Reduce Transactions

```
pairsCount = transactions★.flatMap(lambda transaction: findPairs(transaction, validationSet, candidateSize))
pairsCount = pairsCount.reduceByKey(lambda x, y: x + y)
pairsCount = pairsCount.filter(lambda item: item[1] > min_frequency)
```

Find Pairs

```
def findPairs(transaction:list, validationSet:tuple, candidateSize:int):
    if len(transaction) < candidateSize: return None
    validItems = set(filter(lambda item: item in validationSet, transaction))
    if len(validItems) < candidateSize: return tuple()

    pairs = combinations(transaction, candidateSize) ★
    return tuple((pair, 1) for pair in pairs)
```

How far can I push it?



Why not consolidating the item removal? It will avoid redundant future checks.



Keep calm and just count frequencies.



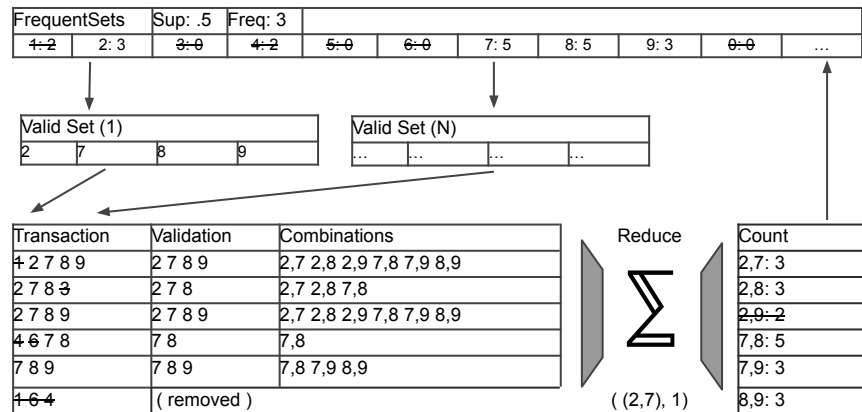
Progressive expansion of N allow to cheaply apply Erosion keeping the future expansions limited.



PCY Erode

Steps

1. Count single transaction items
2. Obtain FrequentSets of size N=1.
3. Make the ValidationSet, the FrequentSets on size N.
4. Remove each line with less then N items.
5. Remove from each transaction line the items not present in the itemsets.
6. Remove each line with less then N item.
7. For each transaction line generate a list of actual sets of size N from the combination items of that row.
8. Each transaction line generates multiple sets, make them flat.
9. Reduce by key and filter.
10. Filter candidates and obtain Frequent Itemsets.



PCY Erode: Focus

Reduce Transactions

```
transactions = reduceTransactions(transactions, validationSet, candidateSize)

def reduceTransactions(transactions: pyspark.RDD, validationSet, candidateSize:int)->pyspark.RDD:
    reducedSet = transactions.map(lambda transaction: filterTransactions(transaction, validationSet, candidateSize))
    reducedSet = reducedSet.filter(lambda transaction: transaction is not None)
    return reducedSet

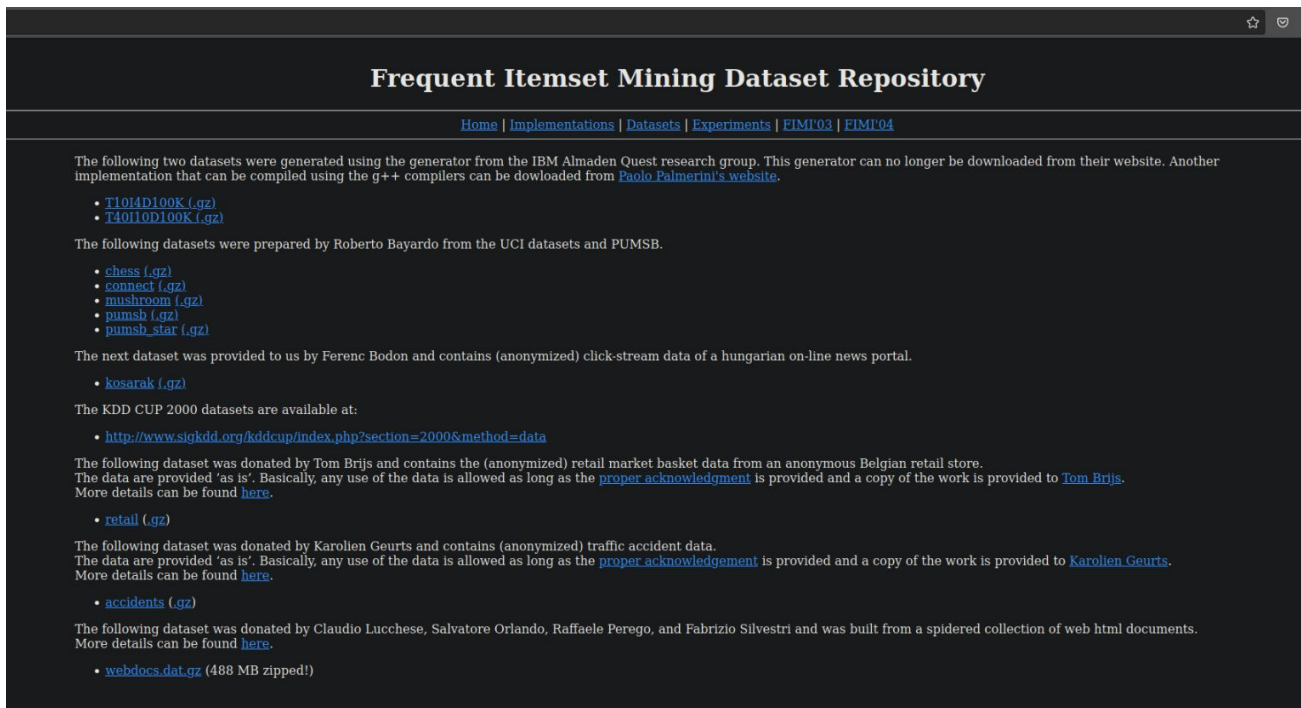
def filterTransactions(transaction, validationSet, candidateSize:int):
    if len(transaction) < candidateSize: return None
    validItems = tuple(filter(lambda item: item in validationSet, transaction))
    if len(validItems) < candidateSize: return None
    return validItems
```

Find Pairs

```
def findPairs(transaction:list|tuple, candidateSize:int) -> tuple:
    pairs = combinations(transaction, candidateSize)
    return tuple((pair, 1) for pair in pairs)
```

DATASETS

Dataset: Overview



The screenshot shows a web browser window displaying the 'Frequent Itemset Mining Dataset Repository' website. The page has a dark theme with a navigation bar at the top containing links: Home, Implementations, Datasets, Experiments, EIMT'03, and EIMT'04. The main content area lists several datasets with descriptions and download links. The datasets listed are T10I4D100K (.gz), T40I10D100K (.gz), chess (.gz), connect (.gz), mushroom (.gz), pumsh (.gz), pumsh_star (.gz), kosarak (.gz), KDD CUP 2000, retail (.gz), accidents (.gz), and webdocs.dat.gz (488 MB zipped!). Each dataset entry includes a brief description of its origin and a link to the dataset file.

Frequent Itemset Mining Dataset Repository

[Home](#) | [Implementations](#) | [Datasets](#) | [Experiments](#) | [EIMT'03](#) | [EIMT'04](#)

The following two datasets were generated using the generator from the IBM Almaden Quest research group. This generator can no longer be downloaded from their website. Another implementation that can be compiled using the g++ compilers can be downloaded from [Paolo Palmerini's website](#).

- [T10I4D100K \(.gz\)](#)
- [T40I10D100K \(.gz\)](#)

The following datasets were prepared by Roberto Bayardo from the UCI datasets and PUMSB.

- [chess \(.gz\)](#)
- [connect \(.gz\)](#)
- [mushroom \(.gz\)](#)
- [pumsh \(.gz\)](#)
- [pumsh_star \(.gz\)](#)

The next dataset was provided to us by Ferenc Bodon and contains (anonymized) click-stream data of a hungarian on-line news portal.

- [kosarak \(.gz\)](#)

The KDD CUP 2000 datasets are available at:

- <http://www.sigkdd.org/kddcup/index.php?section=2000&method=data>

The following dataset was donated by Tom Brijs and contains the (anonymized) retail market basket data from an anonymous Belgian retail store. The data are provided 'as is'. Basically, any use of the data is allowed as long as the [proper acknowledgement](#) is provided and a copy of the work is provided to [Tom Brijs](#). More details can be found [here](#).

- [retail \(.gz\)](#)

The following dataset was donated by Karolien Geurts and contains (anonymized) traffic accident data. The data are provided 'as is'. Basically, any use of the data is allowed as long as the [proper acknowledgement](#) is provided and a copy of the work is provided to [Karolien Geurts](#). More details can be found [here](#).

- [accidents \(.gz\)](#)

The following dataset was donated by Claudio Lucchese, Salvatore Orlando, Raffaele Perego, and Fabrizio Silvestri and was built from a spidered collection of web html documents. More details can be found [here](#).

- [webdocs.dat.gz](#) (488 MB zipped!)

Dataset: Overview

				Support		
Name	Size	Lines	Num Items	AVG	MAX	AVG Items
Mushroom	570.408	8.124	119	0.193	1.000	23
Retail	4.167.490	88.162	16.470	0.001	0.575	10
50K10	2.008.453	50.000	869	0.012	0.078	10
25K40	3.869.620	25.000	942	0.042	0.288	40
100K10	4.022.055	100.000	870	0.0116	0.078	10
100K40	15.478.113	100.000	942	0.042	0.287	40

Dataset: Mushroom

General

Size	570.408
Num lines	8.124
Num Items	119
Support avg	0.193
Support max	1


Frequency

Freq min	4
Freq max	8.124
Freq avg	1.570


Transactions

Items min	23
items max	23
items avg	23

Items frequencies



Items per transaction



Dataset: Retail

General

Size	4.167.490
Num lines	88.162
Num Items	16.470
Support avg	0.001
Support max	0.575


Frequency

Freq min	1
Freq max	50.675
Freq avg	55


Transactions

Items min	1
items max	76
items avg	10

Items frequencies



Items per transaction



Dataset: 50K10

General

Size	2.008.453
Num lines	50.000
Num Items	869
Support avg	0.012
Support max	0.078


Frequency

Freq min	2
Freq max	3.905
Freq avg	581


Transactions

Items min	1
items max	29
items avg	10

Items frequencies



Items per transaction



Dataset: 25K40

General

Size	3.869.620
Num lines	25.000
Num Items	942
Support avg	0.042
Support max	0.288


Frequency

Freq min	1
Freq max	7188
Freq avg	1051


Transactions

Items min	9
items max	72
items avg	40

Items frequencies



Items per transaction



Dataset: 100K10

General

Size	4.022.055
Num lines	100.000
Num Items	870
Support avg	0.012
Support max	0.078


Frequency

Freq min	1
Freq max	7828
Freq avg	1161


Transactions

Items min	1
items max	29
items avg	10

Items frequencies



Items per transaction



Dataset: 100K40

General

Size	15.478.113
Num lines	100.000
Num Items	942
Support avg	0.042
Support max	0.287


Frequency

Freq min	5
Freq max	28738
Freq avg	4204


Transactions

Items min	4
items max	77
items avg	40

Items frequencies



Items per transaction



RESULTS

Results:

	Support	Itemset	apriori		pcy	pcy erode	son
			py	spark			
mushroom	0.9	9	6	13	6	3	8
	0.7	31	18	111	164	4	9
	0.5	153	81	136	1.417	5	16
	0.3	2735	1.423	OOM	24.399	93	140

retail	0.5	1	531	96	3	3	60
	0.3	3	1.577	253	7	3	98
	0.1	9	4.803	723	140	3	221

25K40	0.1	81	1.740	254	9	3	137
100K40	0.1	82	7.115	1.040	26	6	940

50K10	0.01	384	---	---	---	4	426
	0.03	58	---	---	---	2	94
	0.06	4	---	---	---	2	31
retail	0.005	580	---	---	---	7	12.298
	0.05	16	---	---	---	3	339

Results: Cost of a zero

		apriori		pcy	pcy erode	son
		py	spark			
retail	0.7	1	2	3	2	41
	0.9	1	2	2	2	41
100K10	0.1	1	2	2	2	46
	0.3	1	2	2	2	45
	0.5	1	2	3	2	44
	0.7	1	2	2	2	45
	0.9	1	2	2	2	44
100K40	0.3	2	3	3	3	45
	0.5	2	3	3	3	45
	0.7	2	3	3	3	45
	0.9	2	3	3	3	45

		apriori		pcy	pcy erode	son
		py	spark			
50K10	0.1	1	2	2	2	24
	0.3	1	2	2	2	24
	0.3	1	2	2	2	24
	0.7	1	2	2	2	25
	0.9	1	2	2	2	25
100K40	0.3	2	3	3	3	45
	0.5	2	3	3	3	45
	0.7	2	3	3	3	45
	0.9	2	3	3	3	45

Erosion at work: 50K10 @ 0.01 (500)

```
=====
pcy_erode -50K10-0010
=====

Min Freq: 500 = 50000 * 0.01
374
candidateSize: 2
supportSets: ['240', '274', '538', '630', '825', '834', '581', '814', '674', '733', '854', '950', '422', '449', '857', '229', '283', '738', '853', '883', '966', '978', '143', '185', '214', '658', '682',
'782', '947', '970', '227', '390', '192', '208', '279', '280', '496', '530', '675', '720', '914', '932', '217', '161', '175', '490', '571', '623', '960', '130', '461', '862', '900', '147', '411', '572',
'579', '778', '290', '458', '70', '204', '334', '513', '504', '73', '419', '469', '722', '846', '326', '526', '975', '116', '198', '541', '805', '631', '780', '935', '17', '763', '956', '145', '385', '676',
'790', '792', '885', '522', '12', '296', '354', '548', '346', '477', '829', '234', '649', '600', '157', '115', '517', '736', '744', '641', '417', '628', '111', '154', '580', '10', '132', '21', '54', '348',
'100', '48', '319', '112', '140', '285', '387', '93', '583', '122', '718', '1', '69', '797', '110', '509', '611', '995', '33', '336', '598', '470', '992', '897', '259', '45', '162', '378', '716', '8', '413',
'823', '982', '515', '694', '57', '812', '414', '752', '361', '108', '486', '440', '265', '540', '468', '819', '886', '429', '68', '4', '887', '707', '815', '948', '634', '351', '949', '163', '335', '922',
'173', '258', '608', '820', '207', '25', '52', '368', '448', '561', '687', '775', '39', '120', '205', '401', '704', '35', '895', '937', '964', '294', '381', '708', '766', '104', '569', '620', '798', '350',
'529', '809', '71', '597', '618', '183', '276', '653', '706', '878', '177', '424', '795', '910', '125', '392', '27', '78', '921', '803', '266', '523', '614', '888', '944', '43', '480', '874', '151', '890',
'310', '810', '844', '918', '967', '403', '774', '788', '789', '201', '171', '701', '946', '471', '487', '638', '678', '735', '242', '758', '617', '859', '684', '740', '841', '210', '605', '884', '460',
'746', '28', '5', '919', '196', '489', '494', '673', '362', '591', '31', '58', '181', '472', '573', '651', '168', '632', '832', '871', '988', '72', '981', '32', '239', '500', '126', '639', '765', '521',
'594', '606', '236', '952', '90', '593', '941', '423', '516', '6', '913', '577', '343', '527', '989', '97', '574', '793', '427', '37', '55', '275', '51', '534', '906', '576', '373', '665', '963', '349',
'197', '749', '94', '984', '692', '567', '800', '41', '923', '377', '991', '998', '899', '710', '867', '170', '438', '563', '357', '332', '322', '928', '75', '38', '784', '686', '663', '843', '129', '578',
'510', '860', '309', '804', '826', '394', '105', '308', '661', '405', '688', '893', '85', '450', '550', '769', '554', '366']
candidateSize: 3
supportSets: ['227', '390', '217', '346', '722', '825', '829', '682', '39', '704', '789', '368']
candidateSize: 4
supportSets: ['825', '39', '704']
----- RESULTS -----
pcy_erode -50K10-0010
success
384
3.919753313064575
-----
TASK: success 384 3.919753313064575
```

Erosion at work: retail @ 0.005 (440)

```
=====
pcy_erode -retail-0005
=====
Min Freq: 440 = 88162 * 0.005
221
candidateSize: 2
supportSets: ['9', '10', '19', '45', '48', '53', '56', '60', '107', '110', '147', '150', '161', '175', '179', '185', '208', '209', '229', '237', '249', '255', '258', '259', '264', '270', '272', '279', '286',
'301', '334', '338', '345', '365', '389', '408', '413', '426', '441', '449', '464', '488', '490', '522', '548', '570', '571', '581', '623', '649', '664', '675', '677', '703', '772', '783', '790', '806',
'812', '824', '831', '846', '885', '947', '956', '1020', '1043', '1144', '1146', '1198', '1239', '1344', '1404', '1600', '1659', '1693', '1704', '1783', '1814', '1872', '2046', '2080', '2118', '2168', '2383',
'2879', '2958', '4883', '10515', '12929', '15618', '16011', '2', '11', '18', '23', '30', '31', '32', '36', '37', '38', '39', '41', '47', '49', '52', '55', '62', '65', '76', '78', '79', '80', '89', '94',
'101', '103', '105', '117', '123', '136', '155', '156', '170', '178', '186', '201', '225', '242', '251', '260', '261', '267', '269', '271', '281', '297', '310', '340', '371', '374', '379', '396', '398',
'405', '407', '418', '420', '425', '432', '438', '475', '479', '533', '535', '544', '549', '589', '592', '604', '681', '704', '740', '766', '789', '793', '798', '808', '809', '855', '856', '865', '878',
'910', '913', '916', '961', '976', '1004', '1067', '1113', '1135', '1327', '1355', '1393', '1417', '1435', '1479', '1578', '1585', '1677', '1714', '1715', '1867', '1986', '1987', '2051', '2135', '2184',
'2199', '2238', '2284', '2329', '2399', '2437', '3270', '3616', '4336', '8978', '10444', '10446', '12925', '12946', '12959', '13041', '14098', '14933', '15832', '16010', '16217']
candidateSize: 3
supportSets: ['147', '48', '175', '179', '229', '249', '255', '258', '237', '286', '161', '334', '338', '60', '413', '45', '270', '522', '107', '264', '548', '570', '110', '19', '649', '783', '812', '824',
'9', '956', '1146', '677', '301', '1600', '703', '1814', '185', '772', '2958', '10515', '56', '272', '664', '790', '806', '10', '389', '1344', '12929', '16011', '36', '38', '39', '41', '37', '55', '49', '32',
'76', '78', '79', '101', '89', '103', '105', '117', '170', '201', '65', '242', '225', '269', '271', '310', '371', '11', '405', '438', '475', '549', '592', '604', '23', '533', '740', '789', '479', '31', '589',
'766', '123', '1135', '1327', '1393', '1435', '1004', '1578', '544', '18', '976', '2238', '156', '3270', '13041', '12925', '14098', '15832', '16010', '16217']
candidateSize: 4
supportSets: ['147', '48', '179', '249', '255', '258', '237', '286', '60', '270', '548', '110', '824', '9', '1146', '677', '301', '338', '2958', '185', '413', '36', '38', '39', '41', '37', '89', '170', '32',
'225', '105', '371', '65', '475', '271', '310', '2238', '49', '79', '101', '78', '438', '592', '201', '589', '533', '123', '604', '1327', '1393', '13041', '12925', '14098', '16010', '16217']
candidateSize: 5
supportSets: ['286', '48', '110', '36', '38', '39', '41', '170', '32', '89', '65', '225', '310']
candidateSize: 6
supportSets: ['48', '32', '38', '39', '41']
----- RESULTS -----
pcy_erode -retail-0005
success
580
6.335402727127075
-----
TASK: success 580 6.335402727127075
```

[illegible]

Other findings


Hadoop native libraries

```
23/07/02 14:44:47 WARN NativeCodeLoader:
Unable to load native-hadoop library for your platform... using builtin-java classes where applicable.
```



Using Integers

```
def load(sc, data_path, sep=' ', useInt=True) -> pyspark.RDD:
    itemSets = sc.textFile(data_path)
    itemSets = itemSets.map(lambda line: line.strip().split(sep))
    if useInt:
        itemSets = itemSets.map(lambda line: tuple(map(lambda item: int(item), line)))
    return itemSets
```



		Sup	Strings	Integers	Time saved
mushroom	apriori spark	0.7	113	115	-0%
	pcy erode	0.3	93	77	-15%
	apriori python	0.5	81	69	-10%

(pure pyspark)





CONCLUSIONS


Conclusions

- Always plot stats of your dataset as first step, will save tons of time down the line.
- Curse of dimensionality:
 - Join and Cartesian products should be avoided or limited.
 - Alternatives: in-transaction line combinations
 - *Better a Join than a loop: combinations can be an option but can become tricky.*
- Filtering can help reducing the overall computation and if done cheaply doesn't decrease the performance in the worst case scenario.
- Performing an initial filtering even against the Frequent Itemset of size 1, which is anyway computed, can significantly reduce the number items per line, as well as overall lines, greatly decreasing.
- Finding and Verifying Candidate Itemsets can be both wasteful and very expensive.
- In-line combinations when used in combination with counting and filtering can be a scalable alternative to often wasteful cartesian products.
- PCY Erode shows that progressive erosion of the dataset and controlled expansions of the rows, an effective ways to keep under control faster-than-linear growths.

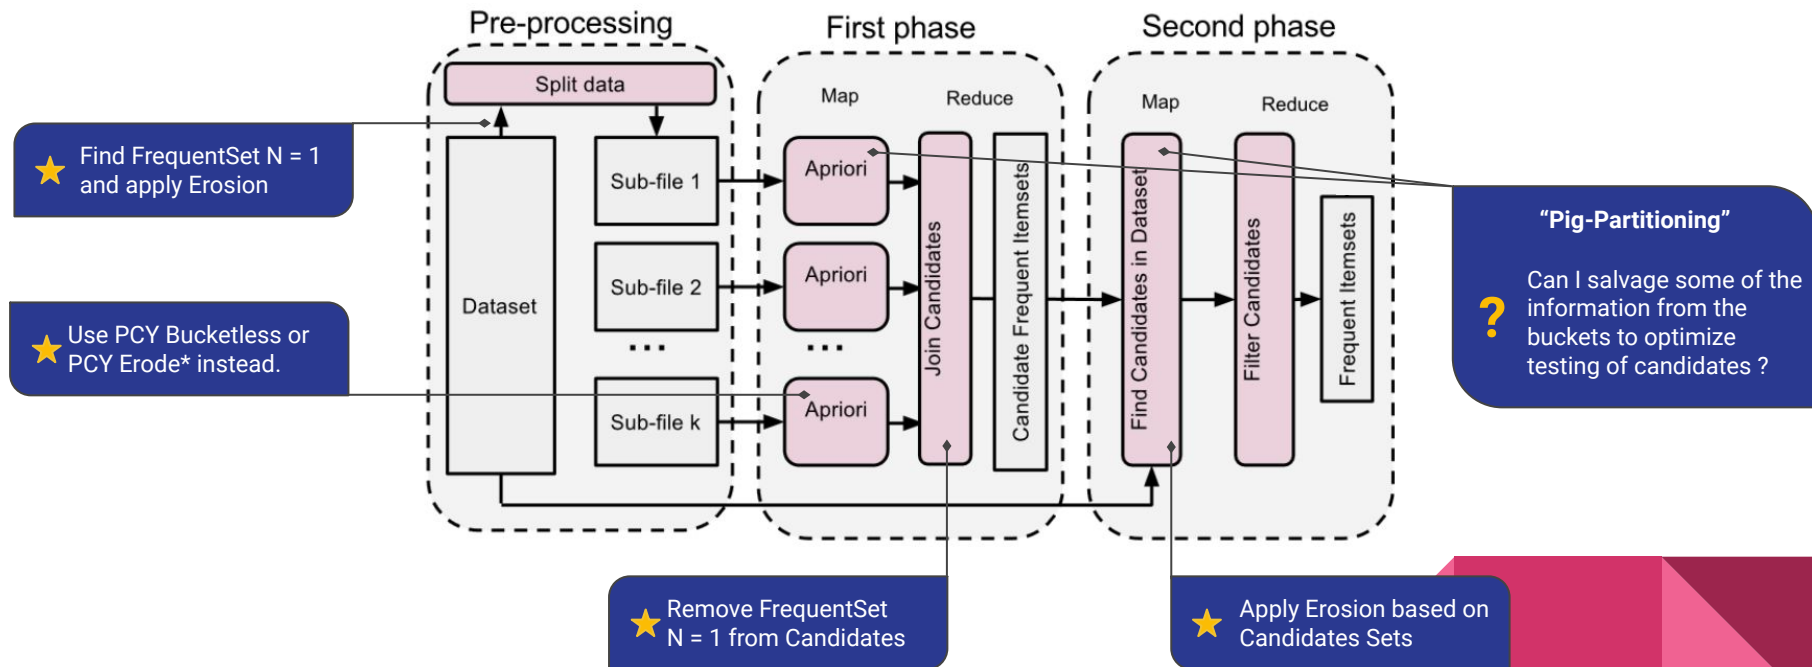
The background is a solid dark blue color. In the top right corner, there is a geometric pattern of triangles in various shades of blue and white, creating a modern, abstract design.

FUTURE

Future

- By the current results I foresee a potential for proving “PCY Erode” with even more precise line filtering and while removing the need of resorting to combinations, potentially reaching an almost guaranteed minimum expansions of each row. However the computational cost of implementing such an aggressive filtering has yet to be accessed and trade-offs has yet to be evaluated.
 - Considering that “PCY Erode” doesn’t use the concept of buckets should be a good algorithm to replace apriori to discover Candidate Itemsets.
 - By the current results I believe it should be possible to apply several optimizations also to the second phase of the SON algorithm by reusing the partitioning generated at the first step, combined with filtering and in-line expansion, to sensibly reduce the size of the cartesian product between the transaction and the candidate itemset.
- 

Future: SON Erode



* need to make a copy of the partition